


SPORES!




Towards Function-Passing Style
in the Age of Concurrency & Distribution

Heather Miller

Scala Days 2014, Berlin, Germany

June 17th, 2014



the fifth annual Scala Workshop

Scala2014

July 28–29th

UPPSALA, SWEDEN

co-located with
ECOOP



The leading forum for *research* & *development* related to the Scala programming language.

AGENDA

SPORES



WHY WE NEED THEM
WHAT THEY ARE

WHAT YOU CAN
DO WITH THEM

¡DEMO!

AGENDA

SPORES

WHY WE NEED THEM
WHAT THEY ARE

WHAT YOU CAN
DO WITH THEM

!DEMO!

(spores in Scala & Javascript)

AGENDA

SPORES

WHY WE NEED THEM
WHAT THEY ARE

WHAT YOU CAN
DO WITH THEM

¡DEMO!

(spores in Scala & Javascript)


SPOILER ALERT:
COOLEST PART OF THE TALK

TWO TRENDS



**DATA-CENTRIC
APPLICATIONS**

**CALLBACKS/REACTIVE.
CLOSURE-HEAVY.**



**FUNCTIONAL
PROGRAMMING**

**APPLY FUNCTIONS TO
IMMUTABLE DATA.**

OBSERVATION:

OBSERVATION:

**THESE TRENDS
ARE COMPLIMENTARY**

OBSERVATION:

**THESE TRENDS
ARE COMPLIMENTARY**

**FUNCTIONAL PROGRAMMING A BOON
TO DATA-CENTRIC PROGRAMMING**

WHY?

①

The basic philosophy to transform immutable data by applying first-class functions.

②

The observation that this functional style simplifies reasoning about data in parallel, concurrent, and distributed code.

HENCE,
THE POPULARITY
OF DATA-PARALLEL
FRAMEWORKS

PARALLEL/ CONCURRENT

Scala's parallel collections

Java 8's monadic
optionally-parallel
collections

Intel's Concurrent Collections

Haskell's Par Monad

DISTRIBUTED

Spark

MapReduce

DATA-PARALLEL FRAMEWORKS

**WHY
ARE SPORES
NECESSARY?**

Well,

**CLOSURES ARE OFTEN A
SOURCE OF HEADACHES**

**YOU CAN'T REALLY
DISTRIBUTE THEM.**

Well,

**CLOSURES ARE OFTEN A
SOURCE OF HEADACHES**

**YOU CAN'T REALLY
DISTRIBUTE THEM.**

**NOT JUST IN SCALA OR JAVA. BUT
CROSS-PARADIGM.**

THE LAUNDRY LIST

PROBLEMS w/ CLOSURES

1. Accidental capture of non-serializable variables (like `this`)
2. Language-specific compilation schemes that create implicit references to objects that are not serializable
3. transitive references that inadvertently hold on to excessively large object graphs creating memory leaks

THE LAUNDRY LIST

PROBLEMS w/ CLOSURES

CONT'D

4. Capturing references to mutable objects, leading to race conditions in a concurrent setting.
5. Unknowingly accessing object members that are not constant such as methods, which in a distributed setting can have logically different meanings on different machines.

motivating example: **SPARK**

```
class MyCoolRddApp {  
  val param = 3.14  
  val log = new Log(...)  
  ...  
  def work(rdd: RDD[Int]) {  
    rdd.map(x => x + param)  
        .reduce(...)  
  }  
}
```

motivating example: **SPARK**

```
class MyCoolRddApp {  
  val param = 3.14  
  val log = new Log(...)  
  ...  
  def work(rdd: RDD[Int]) {  
    rdd.map(x => x + param)  
        .reduce(...)  
  }  
}
```

PROBLEM:

(x => x + param)
not serializable
because it captures
this of type
MyCoolRddApp
which is itself not
serializable

motivating example: **AKKA/FUTURES**

**AKKA ACTOR SPAWNS A FUTURE
TO CONCURRENTLY PROCESS
INCOMING REQS**

```
def receive = {  
  case Request(data) =>  
    future {  
      val result = transform(data)  
      sender ! Response(result)  
    }  
}
```

**NOT A STABLE VALUE!
IT'S A METHOD CALL!**

PROBLEM: Akka actor spawns future to concurrently process incoming results

Scala Improvement Proposal:
<http://docs.scala-lang.org/sips/pending/spores.html>

motivating example: **AKKA/FUTURES**

**AKKA ACTOR SPAWNS A FUTURE
TO CONCURRENTLY PROCESS
INCOMING REQS**

```
def receive = {  
  case Request(data) =>  
    future {  
      val result = transform(data)  
      sender ! Response(result)  
    }  
}
```

**NOT A STABLE VALUE!
IT'S A METHOD CALL!**

PROBLEM: Akka actor spawns future to concurrently process incoming results

Scala Improvement Proposal:
<http://docs.scala-lang.org/sips/pending/spores.html>

ENTER:

SPORES



two types:

①

MAINLINE SPORES

proposed as Scala Improvement Proposal

②

SPORES WITH TYPE CONSTRAINTS

new research published at ECOOP'14

THIS IS ALSO RESEARCH.

Spores: A Type-Based Foundation for Closures in the Age of Concurrency and Distribution

Heather Miller, Philipp Haller¹, and Martin Odersky

EPFL and Typesafe, Inc.¹

{heather.miller, martin.odersky}@epfl.ch and philipp.haller@typesafe.com¹


Abstract. Functional programming (FP) is regularly touted as the way forward for bringing parallel, concurrent, and distributed programming to the mainstream. The popularity of the rationale behind this viewpoint has even led to a number of object-oriented (OO) programming languages outside the Smalltalk tradition adopting functional features such as lambdas and thereby function closures. However,

FOR ALL THE GORY DETAILS...

see our paper accepted for publication at ECOOP'14
<http://infoscience.epfl.ch/record/191239>

ENTER:

SPORES



two types:



MAINLINE SPORES

proposed as Scala Improvement Proposal



SPORES WITH TYPE CONSTRAINTS

new research published at ECOOP'14

mainline SPORES

WHAT ARE THEY?

**SMALL UNITS OF POSSIBLY MOBILE
FUNCTIONAL
BEHAVIOR**

mainline SPORES

WHAT ARE THEY?

A closure-like abstraction for use in distributed or concurrent environments.

GOAL:

Well-behaved closures with controlled environments that can avoid various hazards.

mainline SPORES

POTENTIAL HAZARDS WHEN USING CLOSURES

INCORRECTLY:

- memory leaks
- race conditions due to capturing mutable references
- runtime serialization errors due to unintended capture of references

GOAL:

Well-behaved closures with controlled environments that can avoid various hazards.

WHAT DO SPORES LOOK LIKE?

Basic usage:

```
val s = spore {  
  val h = helper  
  (x: Int) => {  
    val result = x + " " + h.toString  
    println("The result is: " + result)  
  }  
}
```

THE BODY OF A SPORE CONSISTS OF 2 PARTS



a sequence of local value (val) declarations only (the "*spore header*"), and

a closure

A **SPORE** *Guarantees...* (*vs* **CLOSURES**)

1. All captured variables are declared in the spore header, or using capture
2. The initializers of captured variables are executed once, upon creation of the spore
3. References to captured variables do not change during the spore's execution

SPORES & CLOSURES

EVALUATION SEMANTICS:

Remove the spore marker, and the code behaves as before

SPORES & CLOSURES ARE RELATED:

You can write a full function literal and pass it to something that expects a spore.

(Of course, only if the function literal satisfies the spore rules.)

Ok. So. **HOW CAN YOU USE A SPORE?**

IN APIS

If you want parameters to be spores, then you can write it this way

```
def sendOverWire(s: Spore[Int, Int]): Unit = ...  
// ...  
sendOverWire((x: Int) => x * x - 2)
```

Ok. So.

HOW CAN YOU USE A SPORE?

FOR-COMPREHENSIONS

```
def lookup(i: Int): DCollection[Int] = ...
val indices: DCollection[Int] = ...

for { i <- indices
      j <- lookup(i)
    } yield j + capture(i)

trait DCollection[A] {
  def map[B](sp: Spore[A, B]): DCollection[B]
  def flatMap[B](sp: Spore[A, DCollection[B]]): DCollection[B]
}
```


Right,

**WHAT DOES
ALL OF THAT
GET YOU?**

WHAT DOES ALL OF THAT GET YOU?

SINCE...

- Captured expressions are evaluated upon spore creation.

THAT MEANS...

- Spores are like function values with an immutable environment.
- Plus, environment is specified and checked, no accidental capturing.

WHAT DOES ALL OF THAT GET YOU?

OR, GRAPHICALLY...

1
Right after
creation

2
During
execution


SPORES

CLOSURES

Proposed for inclusion in Scala 2.11
<http://docs.scala-lang.org/sips/pending/spores.html>

WHAT DOES ALL OF THAT GET YOU?

OR, GRAPHICALLY...



1
Right after
creation



2
During
execution


SPORES

CLOSURES


Proposed for inclusion in Scala 2.11
<http://docs.scala-lang.org/sips/pending/spores.html>

WHAT DOES ALL OF THAT GET YOU?

OR, GRAPHICALLY...




1
Right after
creation



SPORES

2
During
execution




CLOSURES


Proposed for inclusion in Scala 2.11
<http://docs.scala-lang.org/sips/pending/spores.html>

WHAT DOES ALL OF THAT GET YOU?

OR, GRAPHICALLY...



1
Right after
creation



SPORES

2
During
execution



Proposed for inclusion in Scala 2.11

<http://docs.scala-lang.org/sips/pending/spores.html>

Cool.


Cool. **WHAT IF I
CAPTURE A
SOCKET?**

Cool. WHAT IF I
CAPTURE A
SOCKET?

REALM OF
RESEARCH

ENTER:

SPORES



two types:

①

MAINLINE SPORES

proposed as Scala Improvement Proposal

→ ②

SPORES WITH TYPE CONSTRAINTS

new research published at ECOOP'14

Cool. **WHAT IF I
CAPTURE A
SOCKET?**

**WOULDN'T IT BE NICE IF WE COULD
ADD THESE CONSTRAINTS, IN A
FRIENDLY, AND COMPOSABLE WAY?**

Creating **SPORES** *w/ constraints*

Idea: **KEEP TRACK OF CAPTURED TYPES**
...at compile-time

```
spore { val x: Int = list.size; val a: ActorRef = this.sender  
  (y: Int) => ...  
} exclude[Actor]
```

The spore macro can synthesize precise types
automatically for newly created spores:
(a whitebox macro)

SYNTHESIZED TYPE:

```
Spore[Int, ...] {  
  type Excluded = NoCapture[Actor]  
  type Facts = Captured[Int] with Captured[ActorRef]  
}
```



Composing **SPORES** *w/ constraints*

BASIC COMPOSITION OPERATORS

-  `andThen` *(same as for*
-  `compose` *regular functions)*


How do we synthesize the result type of
`s1 andThen s2`?

RESULT TYPE SYNTHESIZED BY `andThen` **MACRO**

-  type member `Facts` takes "union" of the facts of `s1` and `s2`
-  type member `Excluded`: conjunction of excluded types, needs to check `Facts` to see if possible

Example: Composing **SPORES** w/ constraints

```
val s1: Spore[Int, String] {  
  type Excluded = NoCapture[Actor]  
  type Facts = Captured[Int] with Captured[ActorRef]  
} = ...  
val s2: Spore[String, String] {  
  type Excluded = NoCapture[RDD[Int]]  
  type Facts = Captured[Actor]  
}  
s1 andThen s2 // does not compile
```



Example: Composing **SPORES** *w/ constraints*

```
val s1: Spore[Int, String] {  
  type Excluded = NoCapture[Actor]  
  type Facts = Captured[Int] with Captured[ActorRef]  
} = ...  
val s2: Spore[String, String] {  
  type Excluded = NoCapture[RDD[Int]]  
}  
s1 andThen s2: Spore[Int, String] {  
  type Excluded = NoCapture[Actor] with  
NoCapture[RDD[Int]]  
  type Facts = Captured[Int] with Captured[ActorRef]  
}
```


WHAT DO TYPE CONSTRAINTS BUY US?

- Stronger constraints checked at compile time (not "just" basic spore rules)
- Frameworks can make stronger assumptions about spores created by users.
- Confidence in consuming, creating, and composing spores:
 - * Constraints accumulate monotonically
 - * Constraints are never lost when composing spores
- Less brittleness.

**WHEN
WOULD I
WANT TO
SHIP
A
SPORE?**

HERE ARE **4** EXAMPLES WHICH COULD BE ADVANTAGEOUS (e lots probably more)

- 1** Move functionality to distributed (in-memory) data e.g., Spark
- 2** Shippable stream pipelines e.g., reconfigurable streams
- 3** Hot-swapping actor behavior
- 4** Portable closures e.g., JVM to Javascript

2 SHIPPABLE STREAM PIPELINES

```
Flow(text.split("\\s").toVector).  
  // transform  
  map(line => line.toUpperCase).  
  // print to console (can also use ``foreach(println)``)  
  foreach(transformedLine => println(transformedLine)).  
  onComplete(FlowMaterializer(MaterializerSettings())) {  
    case Success(_) => system.shutdown()  
    case Failure(e) =>  
      println("Failure: " + e.getMessage)  
      system.shutdown()  
  }
```

2 SHIPPABLE STREAM PIPELINES

- Each stage has a closure that deals with incoming streaming data
- However, one could imagine that it could be advantageous that each stage is on the machine that's closest to the data
- Yet we still want the code of the entire pipeline to be assembled on one machine
- That means we have to send pipeline stages together with their closures to different machines after the pipeline has been assembled

**HOW DO
SPORES
HELP?**

2 SHIPPABLE STREAM PIPELINES

✓ Spores ensure that serialization doesn't fail at runtime.

✓ Spores enable different serialization frameworks (e.g., Scala Pickling)

✓ Spores enable restricting types that are captured by each closure.

e.g., if the pipeline is built by an actor, we want to ensure that the enclosing actor is never captured.

2 SHIPPABLE STREAM PIPELINES

```
Flow(text.split("\\s").toVector).  
  // transform  
  map(line => line.toUpperCase).  
  // print to console (can also use ``foreach(println)``)  
  foreach(transformedLine => println(transformedLine)).  
  onComplete(FlowMaterializer(MaterializerSettings())) {  
    case Success(_) => system.shutdown()  
    case Failure(e) =>  
      println("Failure: " + e.getMessage)  
      system.shutdown()  
  }
```

BEFORE

2 SHIPPABLE STREAM PIPELINES

```
Flow(text.split("\\s").toVector).  
  // transform  
  map(spore { line => line.toUpperCase }).  
  // print to console (can also use ``foreach(println)``)  
  foreach(spore { transformedLine => println(transformedLine) }).  
  onComplete(FlowMaterializer(MaterializerSettings())) {  
    case Success(_) => system.shutdown()  
    case Failure(e) =>  
      println("Failure: " + e.getMessage)  
      system.shutdown()  
  }
```

AFTER

3 HOT-SWAPPING ACTOR BEHAVIOR

```
class HotSwapActor extends Actor {  
  import context._  
  
  def receive = {  
    case HotSwap(spore) =>  
      val newBehavior: Receive = { case msg => spore(msg) }  
      become(newBehavior)  
    case ..  
  }  
}
```

4 **PORTABLE** (E.G., BETWEEN JVM & JS) **CLOSURES**

Imagine you have a rich UI on a browser-based client interacting with a server.

If fine-grained information has to be exchanged between client and server, then sending spores can simplify the problem.

- 1** Compose functions based on UI selections.
- 2** Send the composed function, and the server is very simple because it just applies the function.
- 3** No manual translation between low-level message fields to functions applied on the server.

TOTALLY COMPOSABLE. EASILY EXTENDABLE.

4 **PORTABLE** (E.G., BETWEEN JVM & JS) **CLOSURES** **EXAMPLE**

SEARCH TOOL FOR USED CAR OFFERS.

**WEBSITE LETS USERS DEFINE A NUMBER OF
PREFERENCES SUCH AS PRICE RANGE.**

**WHEN ALL PREFERENCES HAVE BEEN SELECTED,
SENT TO SERVER, WHICH FILTERS CARS.**

4 PORTABLE (E.G., BETWEEN JVM & JS) CLOSURES EXAMPLE

SEARCH TOOL FOR USED CAR OFFERS.

ISSUES:

Message containing all user prefs complex.

Extending website with new feature to filter for is complicated, code has to be changed in multiple locations:

- UI code
- encoding pref setting into message to send
- decoding pref setting from message received
- adapting server-side logic for new pref

4 PORTABLE (E.G., BETWEEN JVM & JS) CLOSURES EXAMPLE

CAN DO WITH SPORES IN A WAY WHERE ONLY UI NEEDS TO BE CHANGED!

1. Define spores that filter in code that's shared between client & server.

Each filter spore has type:

`Spore[(Car, Boolean), (Car, Boolean)]`

This allows composing two filters using `andThen`

`val filter = filter1 andThen filter2`

A car matches in the case where

`filter((car, true))._2`

4 PORTABLE (E.G., BETWEEN JVM & JS) CLOSURES EXAMPLE

Example filter spore:

```
def priceRange(from: Int, to: Int) =  
  spore {  
    val localFrom = from  
    val localTo = to  
    (pair: (Car, Boolean)) => {  
      val (car, valid) = pair  
      (car,  
        valid && (car.price >= localFrom && car.price < localTo))  
    }  
  }
```

4 PORTABLE (E.G., BETWEEN JVM & JS) CLOSURES EXAMPLE

2. Compose filters on client side

Suppose there's a collection "selections" which contains filter spores. Then, we simply fold it to get the composed filter:

```
val userPrefs =  
    selections.foldLeft(idFilter)(  
        (f1, f2) => f1 andThen f2  
    )
```

3. Pickle and send the composed spore to the server

4 PORTABLE (E.G., BETWEEN JVM & JS) CLOSURES EXAMPLE

3. On the server side:

Unpickle and use the filter spore to churn through a potentially large dataset.
(Can even use frameworks like Spark for that!)

```
val userPrefs =  
  received.unpickle[Spore[(Car, Boolean), (Car, Boolean)]]  
  
val eligible = carsRdd.filter {  
  car => userPrefs((car, true))._2  
}  
  
// send eligible back to user
```

WHAT'S IN THE RELEASE

Spores implementation as described in SIP-21.

Pickling integration module
(see github.com/scala/pickling)

Support for a subset of type constraints
described in the ECOOP'14 paper.

Alpha version hits sonatype in the next day
or two.

THANK YOU.